

THE AI-NATIVE DOCUMENT STANDARD

JDF the JSON document standard.

Documents that carry their own meaning. Open them in any text editor, render them anywhere, feed them straight to AI. It's just JSON — the way documents should've been all along.

The document is smart. The file is a brick.

You see headings, tables, a signed form. The PDF sees ink smeared at random coordinates. So every tool that opens it has to **reverse-engineer what it meant** — and that's where your tokens, your money and your Friday nights go to die.

Every "AI PDF" tool on the market is just a **duct-tape patch over a file that forgot what it was**. We can do better.

PDF vs JDF

	 PDF	 JDF
Format	Binary blob	Human-readable JSON
Open with	Adobe Reader (~250 MB 💀)	Any text editor or JDF Reader (~14 MB)
Edit a typo	Acrobat / Word roundtrip 😬	Double-click → type → done
Auto-save while editing	nope	✓ debounced, with a status dot
Raw view	lol no	Live two-way JSON view
git diff	Meaningless binary 📄	Clean, line-level
Generate from code	reportlab, pdfkit... good luck	<code>JSON.stringify(doc)</code> . that's it.
Schema validation	none	JSON Schema + Ajv CLI
Search it in code	Custom binary parsing 😬	<code>grep</code> , <code>jq</code> , <code>ripgrep</code>
Vendor lock-in	Adobe 🗝️	zero — it's just JSON 🐦

THE WHOLE THING IN ONE LINE

A document should **carry its own meaning** — not make every reader rebuild it from scratch. 🧠

Giving AI a JDF is handing the builder the **blueprint** 📐. Giving it a PDF is handing over a **blurry photo of the building** 📷 and going "you figure it out."

It already speaks JSON. The friction just disappears.



Actually accurate

"What was Q2 revenue?" → the answer's right there. No scanning, no invented rows.



Cheaper every call

Clean structure = shorter prompts. Stop paying to feed the model garbage.



Nothing goes missing

Tables stay tables, images stay images, signatures stay signed. No [image] tombstones.



AI writes it back

Models read **and** generate JDF natively. Ask for a report, get a real valid doc — not vibes.



Search that slaps

Instant, exact retrieval against real structure — not a hopeful guess at one.



Runs on your box

It's a text file. No vendor, no cloud, no "call sales." Fully yours, fully auditable.

Half your pipeline just got deleted

Pipeline stage	 PDF	 JDF
Parse / extract	pdfplumber / pymupdf — layout analysis, font heuristics, OCR fallback	<code>JSON.parse()</code> — structure's already in the file
Chunking	Token-window splits that slice through tables, lists, footnotes	Every element is a natural chunk — zero config
Metadata	Synthesised after the fact — and usually wrong	First-class on every element: type, heading, page, position
Embedding noise	Repeated headers / footers / page numbers leak in	Header & footer live in their own tree — never in content
Re-index on edit	Re-parse + re-chunk + re-embed the whole PDF	Diff the JSON, re-embed only what changed
Tables	Cells smear across columns; headers collapse	<code>{ headers, rows }</code> — every cell at its real spot
Images / figures	Dropped or stubbed as <code>[image]</code>	Stored with alt text + a stable anchor for vision steps

Write it once. It just works — everywhere, identically.

Same file on your desktop, in a website, and in your CI. No "export as", no conversion tax, no "well it looks different here."



DESKTOP

The Reader

Native app to read, edit, and turn any PDF or Markdown into JDF. Mac, Windows, Linux.



WEB

One tag

Drop `<jdf src>` on any page — renders live, scrollable, searchable, fillable. No plugin.



AUTOMATION

Your pipeline

Convert, validate, feed to AI. Built for agents and CI — zero clicking, ships in a script.

Fill it in the browser. Get back **clean data**.



- 1 Turn any doc into a **fillable form** — text, checkboxes, dropdowns, signatures. The lot.
- 2 User fills it on your site, hits save → you get back **structured data, ready to use**.
- 3 No scraping PDF widget streams, no re-typing into a DB. The answers **are** the data.
- 4 Got legacy PDF forms? They come across with every field and value intact. ↻

WHERE THIS IS GOING 🧙

Documents stop being ~~storage~~. They become **interfaces**.

🎓 Education

Content that's instantly interactive, not just readable.

⚖️ Legal & finance

Contracts and filings you can actually query in real time.

🔌 Every app

One shared structure every tool understands out of the box.

PDF fades out the way fax did once email showed up. 🖨️ → 📧 We're just early.



JSON Document Format

YOUR MOVE 🙌

Ship it. Today.

1

file, everywhere

3

ways to run it

\$0

open source, MIT

∞

opens forever, no vendor

desktop `brew install jdf`

web `<jdf src="doc.jdf"></jdf>`

pipeline `npx @uurtech/jdf-cli convert paper.pdf`

github.com/uurtech/jdf

free & open source ❤️ built by uurtech